

Harsh Seksaria

2048011

2 - MDS

Machine Learning Lab - 4

PCA v/s LDA

27 February 2021

CHRIST (Deemed to be University)

```
In [1]: #Import statements
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: #ML libraries import
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, roc_auc_score
from sklearn.feature_selection import VarianceThreshold
from sklearn.preprocessing import StandardScaler
```

```
In [3]: #Using the inbuilt breast cancer dataset from scikit Learn
from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()
f1 = pd.DataFrame(cancer['data'], columns=cancer['feature_names'])
f1.head()
```

Out[3]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	...	worst radius	worst texture	worst perimeter
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871	...	25.38	17.33	184.6
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667	...	24.99	23.41	158.8
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999	...	23.57	25.53	152.5
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.09744	...	14.91	26.50	98.8
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883	...	22.54	16.67	152.2

5 rows × 30 columns

Data Pre-processing

1. Dataset shape

```
In [4]: print("Number of rows in dataset:", f1.shape[0])
print("Number of columns in dataset:", f1.shape[1])
```

Number of rows in dataset: 569
Number of columns in dataset: 30

2. Null Values

```
In [5]: fl.isnull().sum()
```

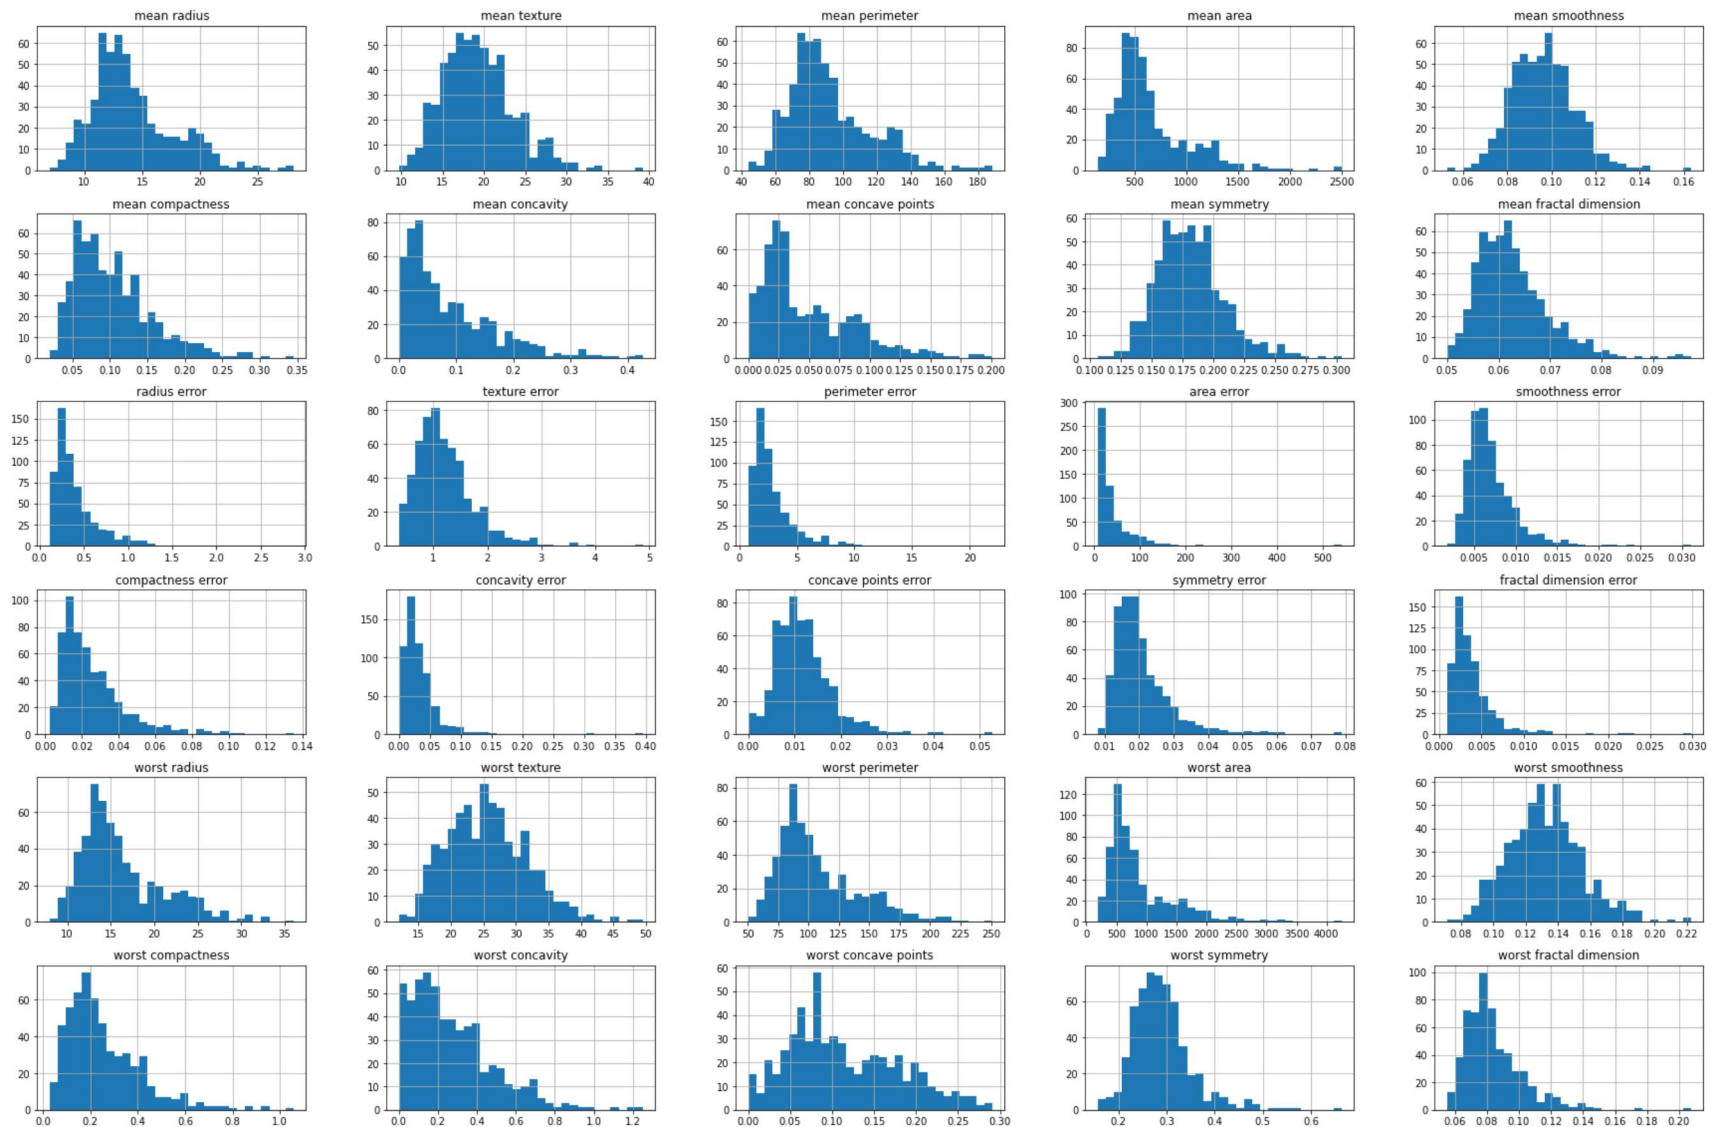
```
Out[5]: mean radius          0
        mean texture        0
        mean perimeter       0
        mean area            0
        mean smoothness      0
        mean compactness     0
        mean concavity        0
        mean concave points  0
        mean symmetry         0
        mean fractal dimension 0
        radius error         0
        texture error        0
        perimeter error       0
        area error           0
        smoothness error     0
        compactness error     0
        concavity error      0
        concave points error  0
        symmetry error        0
        fractal dimension error 0
        worst radius         0
        worst texture        0
        worst perimeter       0
        worst area           0
        worst smoothness     0
        worst compactness    0
        worst concavity      0
        worst concave points  0
        worst symmetry        0
        worst fractal dimension 0
        dtype: int64
```

There are no null values in the dataset.

3. Data Distribution

```
In [6]: fl.hist(bins=30, figsize=(30, 20))
```

```
Out[6]: array([[<AxesSubplot:title={'center':'mean radius'}>,
<AxesSubplot:title={'center':'mean texture'}>,
<AxesSubplot:title={'center':'mean perimeter'}>,
<AxesSubplot:title={'center':'mean area'}>,
<AxesSubplot:title={'center':'mean smoothness'}>],
[<AxesSubplot:title={'center':'mean compactness'}>,
<AxesSubplot:title={'center':'mean concavity'}>,
<AxesSubplot:title={'center':'mean concave points'}>,
<AxesSubplot:title={'center':'mean symmetry'}>,
<AxesSubplot:title={'center':'mean fractal dimension'}>]],
[<AxesSubplot:title={'center':'radius error'}>,
<AxesSubplot:title={'center':'texture error'}>,
<AxesSubplot:title={'center':'perimeter error'}>,
<AxesSubplot:title={'center':'area error'}>,
<AxesSubplot:title={'center':'smoothness error'}>],
[<AxesSubplot:title={'center':'compactness error'}>,
<AxesSubplot:title={'center':'concavity error'}>,
<AxesSubplot:title={'center':'concave points error'}>,
<AxesSubplot:title={'center':'symmetry error'}>,
<AxesSubplot:title={'center':'fractal dimension error'}>]],
[<AxesSubplot:title={'center':'worst radius'}>,
<AxesSubplot:title={'center':'worst texture'}>,
<AxesSubplot:title={'center':'worst perimeter'}>,
<AxesSubplot:title={'center':'worst area'}>,
<AxesSubplot:title={'center':'worst smoothness'}>],
[<AxesSubplot:title={'center':'worst compactness'}>,
<AxesSubplot:title={'center':'worst concavity'}>,
<AxesSubplot:title={'center':'worst concave points'}>,
<AxesSubplot:title={'center':'worst symmetry'}>,
<AxesSubplot:title={'center':'worst fractal dimension'}>]]],
dtype=object)
```



The maximum number of columns in the dataset is not normally distributed. There is skewness in the data.

PCA & LDA

```
In [7]: x = f1  
y = pd.DataFrame(cancer['target'])  
y.columns = ['CancerType']
```

```
In [8]: x.shape, y.shape
```

```
Out[8]: ((569, 30), (569, 1))
```

```
In [9]: #Splitting into training and testing set  
X_train, X_test, Y_train, Y_test = train_test_split(x, y, test_size=0.3, random_state=11, stratify=y)
```

```
In [10]: #Removing constant features using variance threshold  
constant_filter = VarianceThreshold(threshold=0.10)  
  
#Applying the filter on the training set  
constant_filter.fit(X_train)  
  
#Remove constant features from training and test sets  
X_train_filter = constant_filter.transform(X_train)  
X_test_filter = constant_filter.transform(X_test)
```

```
In [11]: #Number of duplicated features in the dataset, We remove one of the feature  
X_train_T = X_train_filter.T  
X_test_T = X_test_filter.T  
X_train_T = pd.DataFrame(X_train_T)  
X_test_T = pd.DataFrame(X_test_T)  
X_train_T.duplicated().sum()
```

```
Out[11]: 0
```

```
In [12]: duplicated_features = X_train_T.duplicated()  
features_to_keep = [not index for index in duplicated_features]  
X_train_unique = X_train_T[features_to_keep].T  
X_test_unique = X_test_T[features_to_keep].T
```

```
In [13]: #standardize the data to get the same scale
scaler = StandardScaler().fit(X_train_unique)
X_train_unique = scaler.transform(X_train_unique)
X_test_unique = scaler.transform(X_test_unique)
X_train_unique = pd.DataFrame(X_train_unique)
X_test_unique = pd.DataFrame(X_test_unique)
X_train_unique.shape, X_test_unique.shape
```

```
Out[13]: ((398, 11), (171, 11))
```

```
In [14]: #Correlation matrix
corrmat = X_train_unique.corr()
corrmat.head()
```

```
Out[14]:
```

	0	1	2	3	4	5	6	7	8	9	10
0	1.000000	0.333688	0.997769	0.986198	-0.104559	0.685085	0.728089	0.969942	0.289167	0.965239	0.941609
1	0.333688	1.000000	0.339211	0.326528	0.370602	0.279599	0.246912	0.362535	0.911823	0.365127	0.345418
2	0.997769	0.339211	1.000000	0.985262	-0.094629	0.703385	0.735965	0.969725	0.295405	0.970500	0.942013
3	0.986198	0.326528	0.985262	1.000000	-0.074554	0.749316	0.804229	0.961981	0.273572	0.958070	0.960621
4	-0.104559	0.370602	-0.094629	-0.074554	1.000000	0.204193	0.078361	-0.119353	0.413891	-0.111507	-0.093683

```
In [15]: #Finding the number of correlated features
def get_correlation(data, threshold):
    corr_col = set()
    corrmat = data.corr()
    for i in range(len(corrmat.columns)):
        for j in range(i):
            if abs(corrmat.iloc[i, j]) > threshold:
                colname = corrmat.columns[i]
                corr_col.add(colname)
    return corr_col

corr_features = get_correlation(X_train_unique, 0.40)
print('correlated features: ', len(set(corr_features)) )
```

```
correlated features: 8
```

```
In [16]: X_train_uncorr = X_train_unique.drop(labels=corr_features, axis = 1)
X_test_uncorr = X_test_unique.drop(labels = corr_features, axis = 1)
X_train_uncorr.shape, X_test_uncorr.shape
```

```
Out[16]: ((398, 3), (171, 3))
```

Dimensionality Reduction using LDA

```
In [17]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
#Transforming the data by using fit_transform()
lda = LDA(n_components=1)
X_train_lda = lda.fit_transform(X_train_uncorr, Y_train)
X_test_lda = lda.transform(X_test_uncorr)
```

D:\Anaconda3\lib\site-packages\sklearn\utils\validation.py:72: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(**kwargs)
```

```
In [18]: #Transformed data
X_train_lda.shape, X_test_lda.shape
```

```
Out[18]: ((398, 1), (171, 1))
```

```
In [19]: from sklearn.ensemble import RandomForestClassifier
def run_randomForest(X_train, X_test, y_train, y_test):
    clf = RandomForestClassifier(n_estimators=100, random_state=0, n_jobs=-1)
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    print('Accuracy on test set: ')
    print(accuracy_score(y_test, y_pred))
```



```
In [20]: %%time
run_randomForest(X_train_lda, X_test_lda, Y_train, Y_test)
```

<ipython-input-19-e37fb2efe457>:4: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
clf.fit(X_train, y_train)
```

Accuracy on test set:

0.8245614035087719

Wall time: 460 ms

```
In [21]: %%time
run_randomForest(X_train, X_test, Y_train, Y_test)
```

<ipython-input-19-e37fb2efe457>:4: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
clf.fit(X_train, y_train)
```

Accuracy on test set:

0.9415204678362573

Wall time: 475 ms

Dimensionality Reduction using PCA

```
In [22]: from sklearn.decomposition import PCA
#Removing the features
pca = PCA(n_components=2, random_state=11)
pca.fit(X_train_uncorr)
PCA(copy=True, iterated_power='auto', n_components=2, random_state=11, svd_solver='auto', tol=0.0, whiten=False)
```

Out[22]: PCA(n_components=2, random_state=11)

```
In [23]: #Training and testing dataset by PCA transformation
X_train_pca = pca.transform(X_train_uncorr)
X_test_pca = pca.transform(X_test_uncorr)
X_train_pca.shape, X_test_pca.shape
```

Out[23]: ((398, 2), (171, 2))

```
In [24]: %%time
run_randomForest(X_train_pca, X_test_pca, Y_train, Y_test)
```

```
<ipython-input-19-e37fb2efe457>:4: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    clf.fit(X_train, y_train)
```

Accuracy on test set:

0.847953216374269

Wall time: 600 ms

```
In [25]: %%time
run_randomForest(X_train, X_test, Y_train, Y_test)
```

```
<ipython-input-19-e37fb2efe457>:4: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    clf.fit(X_train, y_train)
```

Accuracy on test set:

0.9415204678362573

Wall time: 475 ms

```
In [26]: #Checking the accuracy for various selected components
for component in range(1,4):
    pca = PCA(n_components=component, random_state=11)
    pca.fit(X_train_uncorr)
    X_train_pca = pca.transform(X_train_uncorr)
    X_test_pca = pca.transform(X_test_uncorr)
    print('Selected Components: ', component)
    run_randomForest(X_train_pca, X_test_pca, Y_train, Y_test)
    print()
```

Selected Components: 1

<ipython-input-19-e37fb2efe457>:4: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
clf.fit(X_train, y_train)
```

Accuracy on test set:

0.6549707602339181

Selected Components: 2

<ipython-input-19-e37fb2efe457>:4: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
clf.fit(X_train, y_train)
```

Accuracy on test set:

0.847953216374269

Selected Components: 3

<ipython-input-19-e37fb2efe457>:4: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
clf.fit(X_train, y_train)
```

Accuracy on test set:

0.8771929824561403

